

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc.

Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>

Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

```
df = pd.read_csv("/content/drive/MyDrive/ML/Churn_Modelling.csv")
```

Preprocessing.

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	
--------	---------	---------------	-----------	----------------	--

0	2	0.00	1	1	1
1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

df.shape

(10000, 14)

df.describe()

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.000000	1.000000e+04	10000.000000	10000.000000
mean	5000.500000	1.569094e+07	650.528800	38.921800
std	2886.89568	7.193619e+04	96.653299	10.487806
min	1.000000	1.556570e+07	350.000000	18.000000
25%	2500.750000	1.562853e+07	584.000000	32.000000
50%	5000.500000	1.569074e+07	652.000000	37.000000
75%	7500.250000	1.575323e+07	718.000000	44.000000
max	10000.000000	1.581569e+07	850.000000	92.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700

std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

df.isnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	False	False	False	False	False	False
False						
1	False	False	False	False	False	False
False						
2	False	False	False	False	False	False
False						
3	False	False	False	False	False	False
False						
4	False	False	False	False	False	False
False						
...
...						
9995	False	False	False	False	False	False
False						
9996	False	False	False	False	False	False
False						
9997	False	False	False	False	False	False
False						
9998	False	False	False	False	False	False
False						
9999	False	False	False	False	False	False
False						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
...
9995	False	False	False	False	False	False
9996	False	False	False	False	False	False
9997	False	False	False	False	False	False
9998	False	False	False	False	False	False
9999	False	False	False	False	False	False

	EstimatedSalary	Exited
0	False	False
1	False	False

2	False	False
3	False	False
4	False	False
...
9995	False	False
9996	False	False
9997	False	False
9998	False	False
9999	False	False

[10000 rows x 14 columns]

df.isnull().sum()

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	RowNumber	10000 non-null	int64
1	CustomerId	10000 non-null	int64
2	Surname	10000 non-null	object
3	CreditScore	10000 non-null	int64
4	Geography	10000 non-null	object
5	Gender	10000 non-null	object
6	Age	10000 non-null	int64
7	Tenure	10000 non-null	int64
8	Balance	10000 non-null	float64
9	NumOfProducts	10000 non-null	int64
10	HasCrCard	10000 non-null	int64
11	IsActiveMember	10000 non-null	int64
12	EstimatedSalary	10000 non-null	float64
13	Exited	10000 non-null	int64

```
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
df.dtypes
```

```
RowNumber      int64
CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age             int64
Tenure          int64
Balance         float64
NumOfProducts  int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
Exited          int64
dtype: object
```

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
       'Geography',
       'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
       'HasCrCard',
       'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1)
#Dropping the unnecessary columns
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	619	France	Female	42	2	0.00
1	608	Spain	Female	41	1	83807.86
2	502	France	Female	42	8	159660.80
3	699	France	Female	39	1	0.00
4	850	Spain	Female	43	2	125510.82

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1

3	0	0	93826.63	0
4	1	1	79084.10	0

Visualization

```
def visualization(x, y, xlabel):
    plt.figure(figsize=(10,5))
    plt.hist([x, y], color=['red', 'green'], label = ['exit',
    'not_exit'])
    plt.xlabel(xlabel, fontsize=20)
    plt.ylabel("No. of customers", fontsize=20)
    plt.legend()
```

```
df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```


```
visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:3208:
VisibleDeprecationWarning: Creating an ndarray from ragged nested
sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays
with different lengths or shapes) is deprecated. If you meant to do
this, you must specify 'dtype=object' when creating the ndarray.
```

```
    return asarray(a).size
```


```
/usr/local/lib/python3.7/dist-packages/matplotlib/cbook/__init__.py:13
76: VisibleDeprecationWarning: Creating an ndarray from ragged nested
sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays
with different lengths or shapes) is deprecated. If you meant to do
this, you must specify 'dtype=object' when creating the ndarray.
```

```
    X = np.atleast_1d(X.T if isinstance(X, np.ndarray) else
np.asarray(X))
```



```
df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']

visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

```
X =
df[['CreditScore','Gender','Age','Tenure','Balance','NumOfProducts','HasCrCard','IsActiveMember','EstimatedSalary']]
states = pd.get_dummies(df['Geography'],drop_first = True)
gender = pd.get_dummies(df['Gender'],drop_first = True)
```

```
df = pd.concat([df,gender,states], axis = 1)
```

Splitting the training and testing Dataset

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	619	France	Female	42	2	0.00
1	608	Spain	Female	41	1	83807.86
1	502	France	Female	42	8	159660.80
2	699	France	Female	39	1	0.00
3	850	Spain	Female	43	2	125510.82

	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Male	Germany	Spain
0	1	1	101348.88	1	0	0	
0							
1	0	1	112542.58	0	0	0	
1							
2	1	0	113931.57	1	0	0	
0							
3	0	0	93826.63	0	0	0	
0							
4	1	1	79084.10	0	0	0	
1							

```
X =
df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
    'IsActiveMember', 'EstimatedSalary', 'Male', 'Germany', 'Spain']]

y = df['Exited']

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()

X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

X_train
array([[ 0.18602342, -1.03446063, -1.38305456, ..., -1.08316117,
         1.7561361 , -0.57691038],
       [ 0.59018077,  0.87112474,  0.69263926, ..., -1.08316117,
         1.7561361 , -0.57691038],
       [-0.290675   , -0.27222648, -1.72900353, ..., -1.08316117,
        -0.56943195,  1.73337147],
       ...,
       [ 1.27413937,  0.29944913,  1.03858823, ..., -1.08316117,
         1.7561361 , -0.57691038],
       [ 0.99433812, -1.32029844, -0.69115662, ..., -1.08316117,
        -0.56943195,  1.73337147],
       [-1.39951697,  0.10889059,  1.73048617, ...,  0.92322364,
        -0.56943195,  1.73337147]])

X_test
array([[ -0.67410634, -0.17694721,  0.34669029, ...,  0.92322364,
         1.7561361 , -0.57691038],
```



```
[ -0.5186612 , -0.08166794,  1.3845372 , ..., -1.08316117,
  -0.56943195, -0.57691038],
[  0.20674943, -1.32029844,  0.69263926, ...,  0.92322364,
  -0.56943195, -0.57691038],
...,
[  0.94252308,  1.34752109, -0.69115662, ...,  0.92322364,
  -0.56943195,  1.73337147],
[  0.09275633,  0.01361132, -1.38305456, ..., -1.08316117,
  1.7561361 , -0.57691038],
[  0.49691369,  1.25224182, -1.38305456, ..., -1.08316117,
  1.7561361 , -0.57691038]])
```

Building the Classifier Model using Keras

import keras *#Keras is the wrapper on the top of tensorflow*
#Can use Tensorflow as well but won't be able to understand the errors initially.

from keras.models import Sequential *#To create sequential neural network*

from keras.layers import Dense *#To create hidden layers*

classifier = Sequential()

#To add the layers

#Dense helps to construct the neurons

#Input Dimension means we have 11 features

Units is to create the hidden layers

#Uniform helps to distribute the weight uniformly

classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "uniform"))

classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) *#Adding second hidden layers*

classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) *#Final neuron will be having sigmoid function*

classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy']) *#To compile the Artificial Neural Network. Used Binary crossentropy as we just have only two output*

classifier.summary() *#3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neuron in last*

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		

dense (Dense)	(None, 6)	72
dense_1 (Dense)	(None, 6)	42
dense_2 (Dense)	(None, 1)	7

```

=====
Total params: 121
Trainable params: 121
Non-trainable params: 0

```

```

classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the
ANN to training dataset

```

```

Epoch 1/50
700/700 [=====] - 1s 911us/step - loss:
0.4948 - accuracy: 0.7971
Epoch 2/50
700/700 [=====] - 1s 942us/step - loss:
0.4291 - accuracy: 0.7977
Epoch 3/50
700/700 [=====] - 1s 896us/step - loss:
0.4238 - accuracy: 0.8017
Epoch 4/50
700/700 [=====] - 1s 919us/step - loss:
0.4185 - accuracy: 0.8230
Epoch 5/50
700/700 [=====] - 1s 894us/step - loss:
0.4144 - accuracy: 0.8281
Epoch 6/50
700/700 [=====] - 1s 921us/step - loss:
0.4108 - accuracy: 0.8299
Epoch 7/50
700/700 [=====] - 1s 907us/step - loss:
0.4081 - accuracy: 0.8311
Epoch 8/50
700/700 [=====] - 1s 889us/step - loss:
0.4064 - accuracy: 0.8316
Epoch 9/50
700/700 [=====] - 1s 888us/step - loss:
0.4050 - accuracy: 0.8323
Epoch 10/50
700/700 [=====] - 1s 893us/step - loss:
0.4037 - accuracy: 0.8333
Epoch 11/50
700/700 [=====] - 1s 909us/step - loss:
0.4029 - accuracy: 0.8324
Epoch 12/50
700/700 [=====] - 1s 931us/step - loss:
0.4023 - accuracy: 0.8333

```

Epoch 13/50
700/700 [=====] - 1s 877us/step - loss:
0.4012 - accuracy: 0.8347
Epoch 14/50
700/700 [=====] - 1s 910us/step - loss:
0.4011 - accuracy: 0.8324
Epoch 15/50
700/700 [=====] - 1s 917us/step - loss:
0.4001 - accuracy: 0.8344
Epoch 16/50
700/700 [=====] - 1s 939us/step - loss:
0.4002 - accuracy: 0.8329
Epoch 17/50
700/700 [=====] - 1s 881us/step - loss:
0.4000 - accuracy: 0.8354
Epoch 18/50
700/700 [=====] - 1s 902us/step - loss:
0.3995 - accuracy: 0.8346
Epoch 19/50
700/700 [=====] - 1s 901us/step - loss:
0.3992 - accuracy: 0.8337
Epoch 20/50
700/700 [=====] - 1s 903us/step - loss:
0.3991 - accuracy: 0.8349
Epoch 21/50
700/700 [=====] - 1s 904us/step - loss:
0.3988 - accuracy: 0.8320
Epoch 22/50
700/700 [=====] - 1s 881us/step - loss:
0.3986 - accuracy: 0.8346
Epoch 23/50
700/700 [=====] - 1s 878us/step - loss:
0.3985 - accuracy: 0.8341
Epoch 24/50
700/700 [=====] - 1s 878us/step - loss:
0.3978 - accuracy: 0.8346
Epoch 25/50
700/700 [=====] - 1s 886us/step - loss:
0.3982 - accuracy: 0.8347
Epoch 26/50
700/700 [=====] - 1s 920us/step - loss:
0.3973 - accuracy: 0.8359
Epoch 27/50
700/700 [=====] - 1s 895us/step - loss:
0.3977 - accuracy: 0.8350
Epoch 28/50
700/700 [=====] - 1s 933us/step - loss:
0.3976 - accuracy: 0.8350
Epoch 29/50
700/700 [=====] - 1s 895us/step - loss:

0.3974 - accuracy: 0.8350
Epoch 30/50
700/700 [=====] - 1s 890us/step - loss:
0.3970 - accuracy: 0.8359
Epoch 31/50
700/700 [=====] - 1s 903us/step - loss:
0.3972 - accuracy: 0.8360
Epoch 32/50
700/700 [=====] - 1s 905us/step - loss:
0.3972 - accuracy: 0.8357
Epoch 33/50
700/700 [=====] - 1s 887us/step - loss:
0.3970 - accuracy: 0.8351
Epoch 34/50
700/700 [=====] - 1s 906us/step - loss:
0.3970 - accuracy: 0.8361
Epoch 35/50
700/700 [=====] - 1s 969us/step - loss:
0.3967 - accuracy: 0.8346
Epoch 36/50
700/700 [=====] - 1s 914us/step - loss:
0.3963 - accuracy: 0.8360
Epoch 37/50
700/700 [=====] - 1s 900us/step - loss:
0.3965 - accuracy: 0.8360
Epoch 38/50
700/700 [=====] - 1s 905us/step - loss:
0.3964 - accuracy: 0.8359
Epoch 39/50
700/700 [=====] - 1s 916us/step - loss:
0.3964 - accuracy: 0.8349
Epoch 40/50
700/700 [=====] - 1s 918us/step - loss:
0.3960 - accuracy: 0.8366
Epoch 41/50
700/700 [=====] - 1s 890us/step - loss:
0.3957 - accuracy: 0.8366
Epoch 42/50
700/700 [=====] - 1s 914us/step - loss:
0.3962 - accuracy: 0.8360
Epoch 43/50
700/700 [=====] - 1s 882us/step - loss:
0.3955 - accuracy: 0.8361
Epoch 44/50
700/700 [=====] - 1s 870us/step - loss:
0.3953 - accuracy: 0.8369
Epoch 45/50
700/700 [=====] - 1s 1ms/step - loss: 0.3948
- accuracy: 0.8387
Epoch 46/50

```
700/700 [=====] - 1s 1ms/step - loss: 0.3946
- accuracy: 0.8386
Epoch 47/50
700/700 [=====] - 1s 1ms/step - loss: 0.3948
- accuracy: 0.8377
Epoch 48/50
700/700 [=====] - 1s 884us/step - loss:
0.3946 - accuracy: 0.8380
Epoch 49/50
700/700 [=====] - 1s 905us/step - loss:
0.3941 - accuracy: 0.8376
Epoch 50/50
700/700 [=====] - 1s 908us/step - loss:
0.3935 - accuracy: 0.8376
```

```
<keras.callbacks.History at 0x7fd6f5d29410>
```

```
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result

from sklearn.metrics import
confusion_matrix, accuracy_score, classification_report

cm = confusion_matrix(y_test, y_pred)

cm

array([[2308,   71],
       [ 395,  226]])


accuracy = accuracy_score(y_test, y_pred)

accuracy

0.8446666666666667

plt.figure(figsize = (10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

Text(69.0, 0.5, 'Truth')
```



```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.97	0.91	2379
1	0.76	0.36	0.49	621
accuracy			0.84	3000
macro avg	0.81	0.67	0.70	3000
weighted avg	0.83	0.84	0.82	3000